

```
In [89]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [90]: # Read the data
train_df = pd.read_csv('HR-Analytics/train.csv')
test_df = pd.read_csv('HR-Analytics/test.csv')
```

```
In [91]: # Initial exploration
print("Train data shape:", train_df.shape)
print("Test data shape:", test_df.shape)
```

```
Train data shape: (54808, 14)
Test data shape: (23490, 13)
```

```
In [92]: train_df.head()
```

```
Out[92]:
```

	employee_id	department	region	education	gender	recruitment_channel	no_of_trainings	age	previous_year_rating	length_of_serv
0	65438	Sales & Marketing	region_7	Master's & above	f	sourcing	1	35	5.0	
1	65141	Operations	region_22	Bachelor's	m	other	1	30	5.0	
2	7513	Sales & Marketing	region_19	Bachelor's	m	sourcing	1	34	3.0	
3	2542	Sales & Marketing	region_23	Bachelor's	m	other	2	39	1.0	
4	48945	Technology	region_26	Bachelor's	m	other	1	45	3.0	

```
In [93]: train_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 54808 entries, 0 to 54807
Data columns (total 14 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   employee_id                          54808 non-null  int64
1   department                          54808 non-null  object
2   region                             54808 non-null  object
3   education                          52399 non-null  object
4   gender                             54808 non-null  object
5   recruitment_channel                 54808 non-null  object
6   no_of_trainings                    54808 non-null  int64
7   age                               54808 non-null  int64
8   previous_year_rating               50684 non-null  float64
9   length_of_service                  54808 non-null  int64
10  KPIs_met >80%                     54808 non-null  int64
11  awards_won?                       54808 non-null  int64
12  avg_training_score                 54808 non-null  int64
13  is_promoted                       54808 non-null  int64
dtypes: float64(1), int64(8), object(5)
memory usage: 5.9+ MB
```

```
In [94]: train_df.isnull().sum()
```

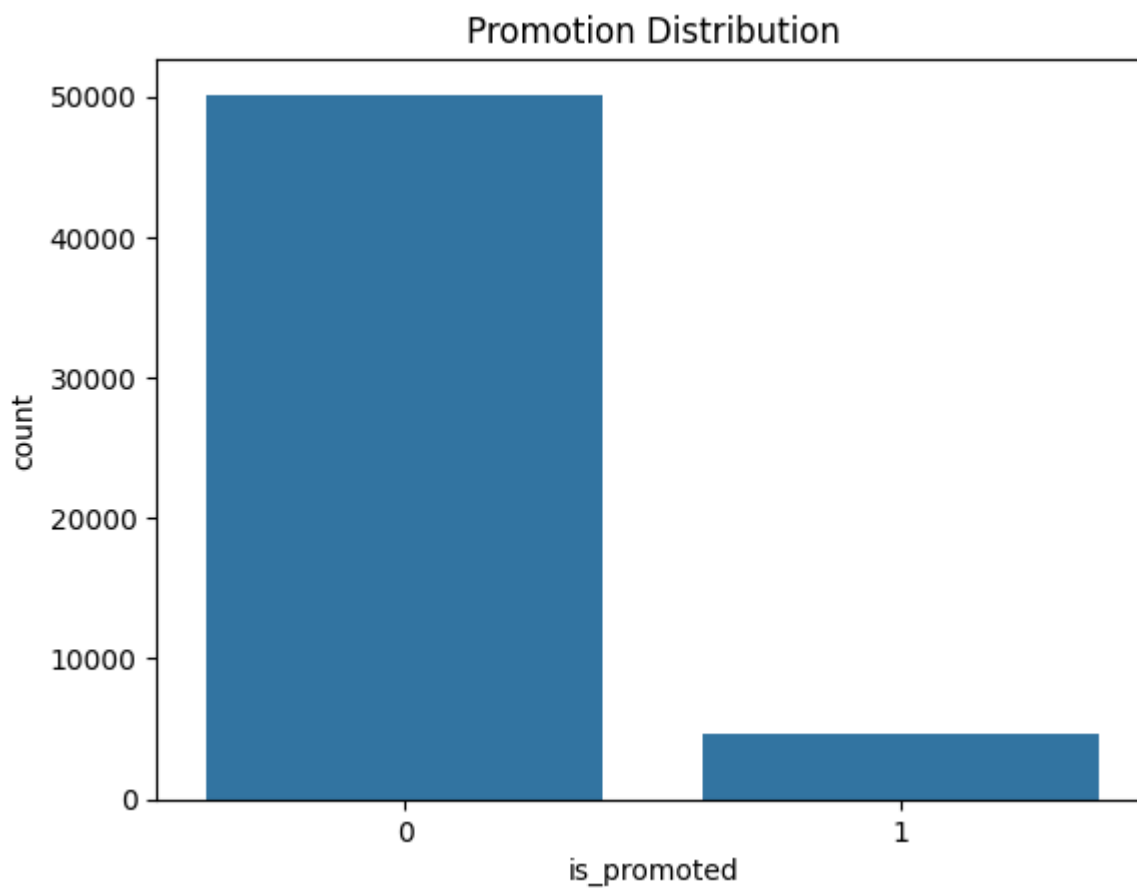
```
Out[94]: employee_id          0
department          0
region              0
education          2409
gender              0
recruitment_channel 0
no_of_trainings     0
age                 0
previous_year_rating 4124
length_of_service   0
KPIs_met >80%       0
awards_won?         0
avg_training_score  0
is_promoted         0
dtype: int64
```

```
In [95]: def fillna(df, column, value):  
        df[column].fillna(value, inplace=True)  
        return df  
  
train_df = fillna(train_df, "education", "unknown")  
train_df = fillna(train_df, "previous_year_rating", 0.0)  
#train_df.drop("employee_id", axis=1, inplace=True)
```


```
In [96]: train_df.isnull().sum()
```

```
Out[96]: employee_id      0  
department      0  
region          0  
education       0  
gender          0  
recruitment_channel  0  
no_of_trainings  0  
age             0  
previous_year_rating  0  
length_of_service  0  
KPIs_met >80%    0  
awards_won?     0  
avg_training_score  0  
is_promoted     0  
dtype: int64
```

```
In [97]: sns.countplot(x='is_promoted', data=train_df)  
plt.title('Promotion Distribution')  
plt.show()
```




```
In [98]: plt.figure(figsize=(8, 6))
sns.countplot(x='department', hue='is_promoted', data=train_df)
plt.title('Promotion by Department')
plt.xticks(rotation=45)
plt.show()
```




```
In [99]: plt.figure(figsize=(6, 6))  
education_counts = train_df['education'].value_counts()
```

```
plt.pie(education_counts, labels=education_counts.index, autopct='%1.1f%%')  
plt.title('Education Level Distribution')  
plt.show()
```



```
In [100... # Select only numeric columns for correlation  
numeric_cols = train_df.select_dtypes(include=np.number).columns  
corr_matrix = train_df[numeric_cols].corr()  
  
# Visualization  
plt.figure(figsize=(16, 12))
```

```
sns.heatmap(corr_matrix, annot=True, fmt=".2f", cmap="coolwarm")  
plt.title("Feature Correlation Matrix (Numeric Columns Only)")  
plt.show()
```



employee

no_of_traini

previous_year_ra

length_of_ser


KPIs_met >=

awards_w

avg_training_s


is_promo

```
In [101... plt.figure(figsize=(8, 4))
plt.hist(train_df['age'], bins=20, edgecolor='black')
plt.title('Age Distribution')
plt.xlabel('Age')
plt.ylabel('Count')
plt.show()
```



```
In [102... plt.figure(figsize=(8, 4))
sns.boxplot(x='is_promoted', y='avg_training_score', data=train_df)
plt.title('Training Score vs Promotion')
```

```
plt.show()
```



```
In [103... train_df.drop(['employee_id', 'region'], axis=1, inplace=True)
```

```
In [104... from sklearn.preprocessing import LabelEncoder
```

```
categorical_cols = ['department', 'education', 'gender', 'recruitment_channel']  
encoders = {}
```

```
for col in categorical_cols:  
    # Get train categories and most frequent value  
    train_categories = train_df[col].astype(str).unique()  
    mode_value = train_df[col].mode()[0]  
  
    # Initialize encoder  
    le = LabelEncoder()
```

```
# Fit on train data
train_encoded = le.fit_transform(train_df[col].astype(str))

# Process test data - replace unseen categories with mode
test_processed = test_df[col].astype(str).apply(
    lambda x: x if x in le.classes_ else mode_value
)

# Transform both datasets
train_df[col] = train_encoded
test_df[col] = le.transform(test_processed)

encoders[col] = le
```

```
In [105... # Example: Create a binary feature for high training score
test_df['high_training_score'] = (test_df['avg_training_score'] > 80).astype(int)
train_df['high_training_score'] = (train_df['avg_training_score'] > 80).astype(int)
```

```
In [106... from sklearn.feature_selection import SelectKBest, f_classif

X = train_df.drop('is_promoted', axis=1)
y = train_df['is_promoted']

# Select top 10 features
selector = SelectKBest(score_func=f_classif, k=10)
X_selected = selector.fit_transform(X, y)
selected_features = X.columns[selector.get_support()]
print("Selected Features (SelectKBest):", selected_features)
```

```
Selected Features (SelectKBest): Index(['education', 'gender', 'no_of_trainings', 'age', 'previous_year_rating',
    'length_of_service', 'KPIs_met >80%', 'awards_won?',
    'avg_training_score', 'high_training_score'],
    dtype='object')
```

```
In [107... from sklearn.feature_selection import RFE
from sklearn.linear_model import LogisticRegression

estimator = LogisticRegression(max_iter=1000)
rfe_selector = RFE(estimator, n_features_to_select=10, step=1)
rfe_selector = rfe_selector.fit(X, y)
```

```
rfe_features = X.columns[rfe_selector.support_]
print("Selected Features (RFE):", rfe_features)
```

```
Selected Features (RFE): Index(['department', 'education', 'no_of_trainings', 'age',
                                'previous_year_rating', 'length_of_service', 'KPIs_met >80%',
                                'awards_won?', 'avg_training_score', 'high_training_score'],
                                dtype='object')
```

```
In [108... from sklearn.ensemble import RandomForestClassifier

rf = RandomForestClassifier(n_estimators=100, random_state=42)
rf.fit(X, y)
importances = rf.feature_importances_
indices = np.argsort(importances)[::-1][:10]
top_features = X.columns[indices]
print("Top Features (Random Forest):", top_features)
```

```
Top Features (Random Forest): Index(['avg_training_score', 'age', 'length_of_service', 'department',
                                      'previous_year_rating', 'KPIs_met >80%', 'recruitment_channel',
                                      'awards_won?', 'no_of_trainings', 'education'],
                                      dtype='object')
```

```
In [109... final_features = list(set(selected_features) | set(rfe_features) | set(top_features))
print("Final Feature Set:", final_features)
```

```
Final Feature Set: ['recruitment_channel', 'length_of_service', 'age', 'no_of_trainings', 'avg_training_score', 'KPI
s_met >80%', 'education', 'previous_year_rating', 'department', 'gender', 'awards_won?', 'high_training_score']
```

```
In [110... from imblearn.over_sampling import SMOTE

X_balanced = train_df[final_features]
y_balanced = train_df['is_promoted']

smote = SMOTE(random_state=42)
X_resampled, y_resampled = smote.fit_resample(X_balanced, y_balanced)
print("Class distribution after SMOTE:", np.bincount(y_resampled))
```

```
Class distribution after SMOTE: [50140 50140]
```

```
In [111... from sklearn.model_selection import train_test_split

X_train, X_val, y_train, y_val = train_test_split(
```

```
X_resampled, y_resampled, test_size=0.2, random_state=42, stratify=y_resampled
)
```

```
In [112... from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report, confusion_matrix, roc_auc_score

lr = LogisticRegression(max_iter=1000, random_state=42)
lr.fit(X_train, y_train)
y_pred_lr = lr.predict(X_val)
y_proba_lr = lr.predict_proba(X_val)[:, 1]

print("Logistic Regression Results:")
print(classification_report(y_val, y_pred_lr))
print("ROC AUC Score:", roc_auc_score(y_val, y_proba_lr))
```

Logistic Regression Results:

	precision	recall	f1-score	support
0	0.71	0.70	0.70	10028
1	0.70	0.72	0.71	10028
accuracy			0.71	20056
macro avg	0.71	0.71	0.71	20056
weighted avg	0.71	0.71	0.71	20056

ROC AUC Score: 0.7911191503837123

```
In [113... from sklearn.tree import DecisionTreeClassifier

dt = DecisionTreeClassifier(random_state=42)
dt.fit(X_train, y_train)
y_pred_dt = dt.predict(X_val)
y_proba_dt = dt.predict_proba(X_val)[:, 1]

print("Decision Tree Results:")
print(classification_report(y_val, y_pred_dt))
print("ROC AUC Score:", roc_auc_score(y_val, y_proba_dt))
```

Decision Tree Results:

	precision	recall	f1-score	support
0	0.95	0.91	0.93	10028
1	0.91	0.95	0.93	10028
accuracy			0.93	20056
macro avg	0.93	0.93	0.93	20056
weighted avg	0.93	0.93	0.93	20056

ROC AUC Score: 0.9301802082211291

```
In [114... from sklearn.ensemble import RandomForestClassifier

rf = RandomForestClassifier(n_estimators=100, random_state=42)
rf.fit(X_train, y_train)
y_pred_rf = rf.predict(X_val)
y_proba_rf = rf.predict_proba(X_val)[:, 1]

print("Random Forest Results:")
print(classification_report(y_val, y_pred_rf))
print("ROC AUC Score:", roc_auc_score(y_val, y_proba_rf))
```

Random Forest Results:

	precision	recall	f1-score	support
0	0.96	0.93	0.95	10028
1	0.94	0.96	0.95	10028
accuracy			0.95	20056
macro avg	0.95	0.95	0.95	20056
weighted avg	0.95	0.95	0.95	20056

ROC AUC Score: 0.9877335035494552

```
In [115... from xgboost import XGBClassifier

xgb = XGBClassifier(use_label_encoder=False, eval_metric='logloss', random_state=42)
xgb.fit(X_train, y_train)
y_pred_xgb = xgb.predict(X_val)
y_proba_xgb = xgb.predict_proba(X_val)[:, 1]
```

```
print("XGBoost Results:")
print(classification_report(y_val, y_pred_xgb))
print("ROC AUC Score:", roc_auc_score(y_val, y_proba_xgb))
```

/media/z-deb/Local_Disk/python_rut/my_project/env/lib/python3.11/site-packages/xgboost/training.py:183: UserWarning:
[14:32:31] WARNING: /workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.

```
bst.update(dtrain, iteration=i, fobj=obj)
XGBoost Results:
```

	precision	recall	f1-score	support
0	0.90	0.89	0.89	10028
1	0.89	0.90	0.90	10028
accuracy			0.90	20056
macro avg	0.90	0.90	0.90	20056
weighted avg	0.90	0.90	0.90	20056

ROC AUC Score: 0.9701113656790902

```
In [116... import pandas as pd

results = {
    'Model': ['Logistic Regression', 'Decision Tree', 'Random Forest', 'XGBoost'],
    'Accuracy': [
        lr.score(X_val, y_val),
        dt.score(X_val, y_val),
        rf.score(X_val, y_val),
        xgb.score(X_val, y_val)
    ],
    'ROC AUC': [
        roc_auc_score(y_val, y_proba_lr),
        roc_auc_score(y_val, y_proba_dt),
        roc_auc_score(y_val, y_proba_rf),
        roc_auc_score(y_val, y_proba_xgb)
    ]
}

results_df = pd.DataFrame(results)
```

```
print(results_df)
```

	Model	Accuracy	ROC AUC
0	Logistic Regression	0.707270	0.791119
1	Decision Tree	0.926855	0.930180
2	Random Forest	0.946600	0.987734
3	XGBoost	0.895692	0.970111

```
In [117... test_df['education'].fillna('unknown', inplace=True)
```

```
In [118... test_df['previous_year_rating'].fillna(0, inplace=True)
```

```
In [119... test_df.head()
```

```
Out[119...
```

	employee_id	department	region	education	gender	recruitment_channel	no_of_trainings	age	previous_year_rating	length_of_serv
0	8724	8	region_26	0	1	2	1	24	0.0	
1	74430	2	region_4	0	0	0	1	31	3.0	
2	72255	7	region_13	0	1	0	1	31	1.0	
3	38562	5	region_2	0	0	0	3	31	2.0	
4	64486	1	region_29	0	1	2	1	30	4.0	

```
In [120... test_df.isnull().sum()
```



```
Out[120...] employee_id      0
           department    0
           region        0
           education      0
           gender         0
           recruitment_channel 0
           no_of_trainings 0
           age            0
           previous_year_rating 0
           length_of_service 0
           KPIs_met >80%    0
           awards_won?     0
           avg_training_score 0
           high_training_score 0
           dtype: int64
```

```
In [121...] from sklearn.preprocessing import LabelEncoder

# Find categorical columns
categorical_cols = test_df.select_dtypes(include=['object', 'category']).columns

# Apply label encoding to each categorical column
le = LabelEncoder()
for col in categorical_cols:
    test_df[col + '_encoded'] = le.fit_transform(test_df[col])
```

```
In [122...] best_model = RandomForestClassifier(n_estimators=100, random_state=42)
best_model.fit(X_resampled, y_resampled) # Using full SMOTE-resampled data
```

```
Out[122...] ▼      RandomForestClassifier ⓘ ?
RandomForestClassifier(random_state=42)
```

```
In [123...] test_predictions = best_model.predict(test_df[final_features])
```

```
In [124...] submission = pd.DataFrame({
    'employee_id': test_df['employee_id'],
    'is_promoted': test_predictions
})
```

```
submission.to_csv('submission.csv', index=False)
```

```
In [125... import os

# Create output directory if needed
os.makedirs('output', exist_ok=True)

# Create submission DataFrame
submission = pd.DataFrame({
    'employee_id': test_df['employee_id'],
    'is_promoted': test_predictions
})


# Save with explicit path and verification
filepath = os.path.join('output', 'submission.csv')
submission.to_csv(filepath, index=False)

# Verify creation
if os.path.exists(filepath):
    print(f"✅ File saved successfully: {filepath}")
    print(f"📄 File contains {len(submission)} rows")
else:
    print("❌ File creation failed - check permissions/paths")

✅ File saved successfully: output/submission.csv
📄 File contains 23490 rows
```

```
In [126... promotions= pd.read_csv("submission.csv")
```

```
In [127... sns.countplot(x='is_promoted', data=promotions)
plt.title('Promotion Distribution')
plt.show()
```



```
In [128... counts = promotions['is_promoted'].value_counts()  
print(counts)
```

```
is_promoted  
0      21352  
1       2138  
Name: count, dtype: int64
```

```
In [ ]:
```